

Study of various algorithms for identifying tonalities in music pieces, proposal of an alternative approach and its implementation in Clojure

Antoine Passemiers

Université Libre de Bruxelles
apassemi@ulb.ac.be

Abstract

The project involves the discussion of different approaches to automated tone detection in audio files, the prototyping of these methods, a reflexion on their consistency with functional programming, and the implementation of the most appropriate one in Clojure. State of the art comprises various maximum key-profile correlation (MKC) approaches and machine learning algorithms, and this report aims to implement all of them in Python to be able to discuss them according to their accuracy. Also, the second goal is to propose an efficient and custom Clojure implementation.

1. Introduction

The tone of a musical work is characterised by the set of sounds coming from a same diatonic scale. Unlike scales, where musical notes follow each other in a contiguous manner, tones assemble notes that can be disjoined or even overlapped. (Danhauser, 1929) Therefore, our central concern is the analysis of polyphonic melodies, which constitute the overwhelming majority of the modern musical works.

More specifically, we will give consideration to two types of tonality prediction methods : those based on tone profiles, and those based on machine learning algorithms. The former aims to integrate the acquired knowledge on the cognitive theory, whereas the latter use statistical inference to extract the relevant information from the music tracks. The cognitive approach used for identifying tones is partly dependent on the solution design proposed by Ibrahim Sha'ath for the software KeyFinder.

As regards the machine learning techniques, those presented from sections 2.3.2 to 2.3.4 are mainly linked to hidden Markov models and neural networks. The cross-validation of these models is achieved by splitting the database in a training set and a validation set. Being intended to remain brief, the present report focuses on the methodology used, and short descriptions of the

algorithms discussed, with emphasis on how the solution presented differs from these algorithms.

2. Theoretical considerations

2.1. Preprocessing

The audio signal is first extracted from the wav file, and then the average of its channels is calculated if the file has been written in stereo. It is not required to take the pan into account given that mixing channels or analysing the channels separately makes no difference in the spectral domain. Also, there is no indispensable need to consider the whole frequential spectrum given that notes are mainly characterised by their fundamental frequency and that the harmonics of the highest notes are not relevant for estimating chroma vectors. The sampling rate is de facto resampled at a tenth of the standard sampling frequency (4410 Hz). Although this substantial downsampling is likely to create aliasing effects (which can be simply avoided by applying a low-pass filter), no loss of accuracy has been pointed out during the evaluation phase.

2.2. Spectral estimation

There are two broad categories of spectral density estimation techniques : the parametric ones and the non-parametric ones. (H.Hayes, 1996) The techniques of which we will be speaking are all non-parametric algorithms. More specifically, we are going to discuss about methods based on the discrete-time Fourier transform (which are the most often used) and least-squares spectral estimation techniques. The latter have achieved comparable results, both in terms of speed and accuracy.

2.2.1. Direct Spectral Kernel Transform (DSKT) This algorithm has been proposed by Ibrahim Sha'ath, and appears to be a fast and efficient approximation of the Constant-Q Transform implementation as described by Brown and Puckette. The data samples are first multiplied by a Blackman window (to avoid spectral distortions). Then

the spectrum is approximated by the Fast Fourier Transform. The spectral kernels are computed a posteriori by applying cosine windows in the spectral domain, centered on the frequencies of interest. The advantage is that each window can be pre-calculated and reused each time the DSKT needs to be applied. The particularity of the DSKT is that the bandwidth around each frequency of interest is constant, and determined by the parameter Q . To calculate the bounds of a kernel of central frequency f_k , we refer to the following formulas :

$$l_k = (1 - \frac{Q}{2}) \cdot \frac{f_k \cdot N}{s} \quad (1)$$

$$r_k = (1 + \frac{Q}{2}) \cdot \frac{f_k \cdot N}{s} \quad (2)$$

as established by Sha'ath, where N is the resolution of the Fourier transform, and s is the sampling rate (after downsampling).

2.2.2. Vaníček least-squares spectral estimation It is also possible to estimate the spectrum using dedicated regression techniques. We will give focus on the most popular of them : the Vaníček method and the Lomb-Scargle method. A spectrum approximated by one of the latter is called a periodogram. Handling missing data samples and non-uniform data is a rationale for the Vaníček method, since it has been developed in the fields of geology and astrophysics, where data are not received at regular intervals. Furthermore, unlike the Fast Fourier Transform which has a variable spectral resolution (depending on the size of the observation window), the least-squares methods output periodograms of a size that is equal to the number of trial frequencies. Theoretically, this is an interesting feature since the window size can be adjusted to the duration of a beat. This can not be achieved using FFT-based algorithms. The only constraint is that the number of data samples must be higher than the number of trial frequencies.

$$\hat{c} = (\phi^T \phi)^{-1} \phi^T f \quad (3)$$

According to the Vaníček method, the signal is assumed to be a linear combination of sine and cosine waves, with given frequencies of interest. The coefficients of the linear combination are stored in a vector c , and the signal is assumed to be $f \approx \phi \cdot c$. Each column of the matrix ϕ contains the data samples of the sine and cosine waves of a particular trial frequency. Finally, the coefficients are given by the equation 3.

(Stoica et al., 2009)

This method is very similar to the matching pursuit algorithm. The idea of the matching pursuit is to represent the signal as a finite linear combination of given functions, referred as atoms. At iteration k , the atom that minimizes the least squares residuals is weighted and then subtracted from the remainder of the signal. Once every trial frequencies has been treated, the weights constitute the spectral estimation vector.

2.2.3. Lomb-Scargle least-squares spectral estimation As we can see from formula 3, the Vaníček's strategy is lacking consideration for the phases of the trial sine waves. The consequence is that sinusoids of a same frequency can be highly correlated. The new approach requires that each pair of sinusoids at a given frequency is mutually orthogonal at each time t_j . Let's set τ the phase change to apply to both the sine wave and the cosine wave to make them mutually orthogonal for each data sample. Luckily, this phase change can be pre-calculated and is given by the equation 4 (Lomb, 1975).

$$\tan 2\pi f \tau = \frac{\sum_{j=1} \sin 2\pi f t_j}{\sum_{j=1} \cos 2\pi f t_j} \quad (4)$$

Contrary to the formula 3, which consider all the trial frequencies at once, Lomb's approach suggests to consider only one trial frequency at a time, and consider the rest of the spectrum as being random noise. This process is repeated progressively for each trial frequency. The weight $R(f)$ of each trial frequency is then grouped together with the others, resulting in a periodogram. $R(f)$ is a linear combination of the signal-cosine wave correlation and the signal-sine wave correlation, and results in the following equation:

$$\Delta R(f) = \frac{(YC)^2}{CC} + \frac{(YS)^2}{SS} \quad (5)$$

Also, the power of the cosine wave and the sine wave are different for the same trial frequency. To rebalance the previous equation, the terms YC and YS are divided respectively by the power of the cosine wave (CC) and the power of the sine wave (SS) Fortunately, CC and SS can be pre-calculated as well, since the trial frequencies don't change from a music track to another.

$$CC = \sum_{j=1} \cos^2 2\pi f(t_j - \tau) \quad (6)$$

$$SS = \sum_{j=1} \sin^2 2\pi f(t_j - \tau) \quad (7)$$

[[Les valeurs YC et YS cherchent à mesurer respectivement les niveau d'orthogonalité du signal y avec une sinusoïde déphasée de $\pi/2$ et avec une sinusoïde de phase nulle. Ces sinusoïdes peuvent également être pré-calculées. Les formules permettant de calculer YC et YS sont :]]

$$YC = \sum_{j=1} y_j \cos 2\pi f(t_j - \tau) \quad (8)$$

$$YS = \sum_{j=1} y_j \sin 2\pi f(t_j - \tau) \quad (9)$$

Therefore, the only remaining operations that cannot be processed at compile-time are the scalar products of the sinusoids and the current chunk, which reduces drastically the number of required operations to estimate the spectrum.

Contrarily to Fourier-based estimators, the periodogram seems to suffer from spectral imbalance : this phenomenon is illustrated in the appendix B. When the periodogram consists of several peaks, the highest peak will always be relatively much higher than the others, compared to what we can get with the Fourier transform. This effect tend to diminish when the signal

This report proposes a simple solution to this issue. The idea consists in enriching the spectrum with musical notes located in a different chunk of data samples. In fact, an heptatonic musical scale is characterized by a set of seven notes, but it is not common to meet them all at the same moment. Because it is not possible to make a spectral analysis of contiguous notes, the alternative is to overlay them.

2.3. Key identification

Once the spectrum is approximated, it is reshaped into a chroma vector of twelve coefficients. In the framework of the algorithms assessment, it has been found that 6 octaves are sufficient to evaluate the spectrum : the addition of a seventh octave does not improve the accuracy of the predictions. With this knowledge, the spectrum - which has a resolution of 72 coefficients is reshaped into a matrix $M_{i,j}$ of dimension (6×12) . Finally, the chroma vector is obtained by formalising the method developed by Sha'ath :

$$C_i = (1 - p) * \max_j M_{i,j} + p * \sum_j M_{i,j} \quad (10)$$




Figure 1: Profils de tonalités dérivés de l'expérience de Krumhansl, et représentant respectivement la gamme do majeur et la gamme do mineur.




Figure 2: Meilleurs profils de tonalités, déterminés empiriquement

where p is a parameter determined during the evaluation / validation of the whole system. The last step consists in identifying the tonality locally, on the basis of this single chroma vector. We are going to discuss about different methods, the first one relies on the Maximum Key-Profile Correlation (MKC), while the others are based on machine learning algorithms.

2.3.1. Modèle cognitif Cette solution repose sur les expériences menées par Carol L. Krumhansl et Lola L. Cuddy sur la perception des tonalités chez l'Homme. Des gammes incomplètes ont été jouées, suivies de notes additionnelles que les sujets devaient évaluer sur leur capacité à compléter ces gammes. Les scores ont été moyennés pour donner les deux graphiques suivants : TODO : Sha'ath (Sha' ath, 2011) and (Pauws, 2004) (Takeuchi, 1994)

Bien que ces deux profils de tonalités fournissent des résultats satisfaisants lorsque le spectre est estimé à l'aide de la CQT, il ne sont pas adaptés lorsque la méthode de Lomb-Scargle est utilisée, car il a été observé que les coefficients renvoyés par cette dernière sont globalement plus équidistribués. Cette distribution doit se refléter dans les profils de tonalité, c'est pourquoi l'algorithme a été évalué plusieurs fois jusqu'à obtenir les profils fournissant les meilleurs résultats, et donnés ci-dessous :

2.3.2. Hidden Markov Models (HMM) Hidden Markov Models are discrete state machines that

represent multivariate series by their distributions and by the transition probabilities between their states. Their states are called hidden because they are not observed directly, but can be inferred from the observation sequence using a parameter estimation method (such as the expectation-maximization algorithm or the Viterbi algorithm). Each hidden state distinguishes itself by its emission probabilities (or emission probability density function). Unlike more popular machine learning models such as neural networks or support vector machines (SVM), Hidden Markov Models are capable of processing input sequences of variable length. This characteristic is substantial given that music tracks are of variable length by nature.

The old-fashioned way to classify keys using HMMs consists of training a HMM for each distinct music key, and to predict the most probable key of an unlabelled music track by looking for the HMM that maximizes the track's log-likelihood.

TODO : (Peeters, 2006) (Ramage, 2007)

2.3.3. Feedforward neural networks Neural networks have already delivered good results for classifying chords inside a same piece of music. (Osmalskyj et al., 2012)

2.3.4. Input-Output Hidden Markov Models (IO-HMM) The main deficiency of the standard HMM is that it has no discriminative capabilities : if one of the keys appears to be more frequent in the database or if it seems to be easily explained by the HMM, the latter will present high risks of predicting the same key every time, regardless of the input sequence. On the other hand, feedforward neural networks have the defect that they are not suited for temporal data, since they must process feature vectors of fixed size. [[TODO : recurrent networks]]

The Input-Output Hidden Markov Model is well suited for our task because it is affected by none of these deficiencies.

TODO : (Bengio and Frasconi, 1996)

3. Methodology

3.1. Clojure implementation

Clojure is a dynamic and functional programming language based on the Java Virtual Machine, suited for quickly designing multithreaded programs, and featuring a full range of immutable and/or persistent data structures. It benefits from both the advantages of compiled languages and dynamic typed languages, which allows us to write optimized code in a small amount of time, using a fancy syntax. However, although Java libraries can be imported into the

Clojure environment, there is a comparatively small number of libraries that are natively implemented in Clojure. In our case, the majority of the functionalities are implemented exclusively for the project. Because the Clojure machine learning libraries are at their first steps, we are not using any machine learning technique in the design of the final algorithms.

The objective of this section is to describe the functional solution design for Clojure, and compare the Lomb-Scargle periodogram and the direct spectral kernel transform in terms of speed and accuracy. Finally, we will insist on the choices that differ from state of the art from a theoretical point of view.

In order to test the numerous features of the project, unit tests have been implemented. In addition, Tufte (a Clojure profiling library) has been used to assess the performance of the various parts of the system. For all the performance-critical areas of the code, we will see the different measures shown by Tufte.

3.2. Solution design

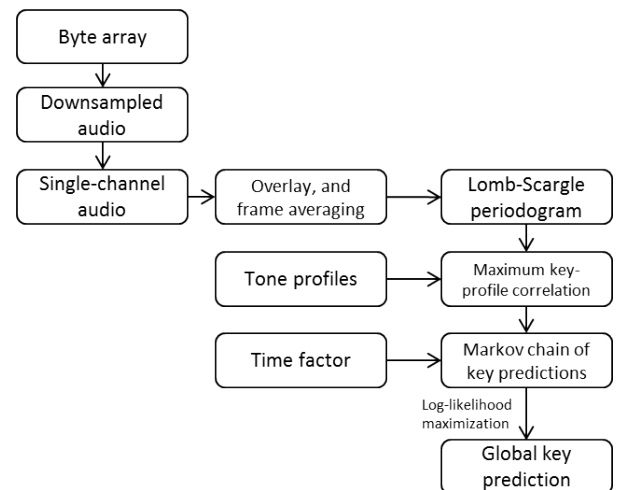


Figure 3: Flowchart : solution design

3.2.1. Downsampling and channel averaging

Since downsampling without filtering has no significative impact on the accuracy, a dedicated wav file reader has been designed to convert a byte stream directly at the target sampling rate. This is to process far less data samples, thus helping to gain speed. In addition, the downsampled signal is then returned as a lazy Clojure sequence, ensuring that we don't evaluate data samples we are not using in the end. The next step consists of, in the case of a multi-channel audio, averaging the

multiple channels. The result of this operation yields a lazy sequence, following the same logic.

3.2.2. Lomb-Scargle periodogram Considering the spectral estimation remains the performance-critical area of the code, regardless of the technique used, type hints have been added to the functions and var names involved. This is to help the compiler infer the type of return values at compile-time. At this stage of the code, we are only handling primitives and *clojure.lang.LazySeq* objects (Clojure lazy sequences).

We know from equations 6 and 7 that the terms CC and SS can be pre-calculated. As a consequence, they are computed at compile-time, for every trial frequency, and can be reused for each spectral estimation. For this reason, a vector of *PreprocessedWaveforms* objects has been statically defined. A *PreprocessedWaveforms* type regroups a specific trial frequency, the associated phase change, sine wave and cosine wave, and the corresponding CC and SS terms.

3.2.3. Direct spectral kernel transform TODO

3.2.4. Datasets The database consists of 415 music tracks recorded in wav format, 44100 Hz stereo, associated with a labelled key (determined beforehand). All keys have been determined by experts and have been made available by Ibrahim Shat'ath as part of its work. For evaluating algorithms based on MKC, all the tracks have been used to tune the parameters to their best value. The latter consist of : the major and minor tone profile coefficients, the size of the sliding window, and the frequency range of the spectrum. For evaluating machine learning techniques, the dataset is split into a training set of 138 music tracks and a validation set of 277 tracks.

3.2.5. Evaluation In the framework of this study, we consider only the twelve heptatonic major scales and the twelve heptatonic minor scales. Two different metrics are used to evaluate the accuracy of the algorithms : the raw accuracy (the ratio between the number of correct predictions to the number of analyzed files) and the MIREX index. The MIREX metric is a weighted sum of the number of correct predictions (C), the number of predictions that are out by a fifth (F), out by a fourth (U), the number of predicted scales that are parallel to the correct ones (P), and the number of predicted scales that are relative to the correct ones (R). The index is then given by :

$$MIREX = (C + \frac{1}{2}.F + \frac{1}{2}.U + \frac{3}{10}.R + \frac{1}{5}.P)/TOTAL$$

Method	Accuracy	MIREX	τ
DSKT (Java FFT)	-	-	-
DSKT (Clojure FFT)	-	-	-
Lomb-Scargle (Clojure)	-	-	-

Table 1: Accuracy assessment, according to the raw accuracy and the MIREX index

Method	Accuracy	MIREX
CQT + MKC	30,7%	-
Lomb-Scargle + MKC	36,6%	49,4%
CQT + HMM	-	-
FFT + IO-HMM	-	-
CQT + IO-HMM	-	-
Lomb-Scargle + IO-HMM	-	-

Table 2: Accuracy assessment, according to the raw accuracy and the MIREX index

4. Results

References

1. Bengio, Y. and Frasconi, P. (1996). Input-output hmm's for sequence processing. *IEEE transactions on neural networks*, vol. 7, no. 5.
2. Danhauser, A. (1929). *Théorie de la Musique*.
3. H.Hayes, M. (1996). *Statistical digital signal processing and modeling*.
4. Lomb, N. R. (1975). Least-squares frequency analysis of unequally spaced data.
5. Osmalskyj, J., Embrechts, J.-J., Piérard, S., and Droogenbroeck, M. V. (9-11 mai 2012). Neural networks for musical chords recognition. *Actes des Journées d'Informatique Musicale (JIM 2012), Mons, Belgique*.
6. Pauws, S. (2004). Musical key extraction from audio. *Proceedings of the 9th International Conference on Digital Audio Effects, DAFx-06, Montreal, Canada*, pages 96–99.
7. Peeters, J. (2006). Musical key estimation of audio signal based on hidden markov modeling of chroma vectors. *Proceedings of the 9th International Conference on Digital Audio Effects, DAFx-06, Montreal, Canada*.

Lomb-Scargle + MKC	Accuracy	MIREX
4096 samples	30,7%	-
8192 samples	37,1%	48,9%
12288 samples	35,9%	48,0%
16384 samples	35,8%	48,9%

Table 3: Scores obtained using Lomb-Scargle + MKC, based on the temporal window size

8. Ramage, D. (2007). Hidden markov models fundamentals. *CS229 Section Notes*.
9. Sha' ath, I. (2011). Estimation of key in digital music recordings. pages 1–64.
10. Stoica, P., Li, J., and He, H. (2009). Spectral analysis of nonuniformly sampled data : A new approach versus the periodogram. *IEEE Transactions on signal processing*, vol 57, no. 3, pages 843–857.
11. Takeuchi, A. H. (1994). Maximum key-profile correlation (mkc) as a measure of tonal structure in music. *Perception & Psychophysics*, pages 335–346.

5. Appendices

5.1. Appendix A - more results

5.2. Appendix B - plots




Figure 4: Dreadlock holiday - 10cc, frame located at




Figure 5: Dreadlock holiday - 10cc, frame located at